

Interprocess Communication (IPC)

Computer Operating Systems
BLG 312E

2016-2017 Spring

Types of Interaction

- three types of interaction between concurrent processes
 - resource sharing
 - communication
 - synchronization

Levels of Interaction

- interaction between processes can be on three levels
 - processes are not aware of each other (competing): using system resources (moderated by operating system)
 - processes are indirectly aware of each other (sharing): resource sharing through mutual exclusion and synchronization
 - processes are directly aware of each other (communicating)

Resource Sharing

- mutual exclusion
 - two types of resources
 - 1) can be used by more than one process at a time (e.g. reading from a file)
 - 2) can be used by only one process at a time
 - due to physical constraints (e.g. some I/O units)
 - if the actions of one process interferes with those of another (e.g. writing to a shared memory location)
- synchronization
 - a process should proceed after another process completes some actions

Example

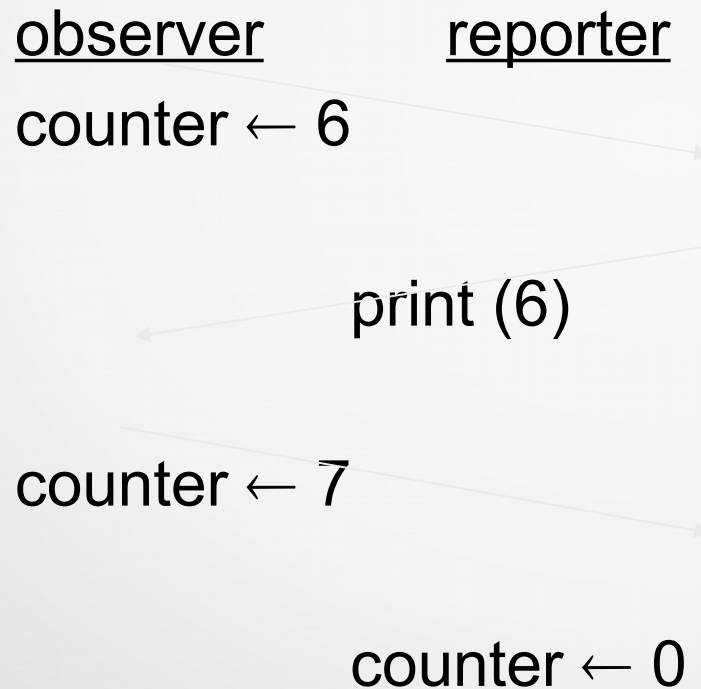
2 processes: Observer and Reporter

counter shared variable

```
observer:
while TRUE {
    observe;
    counter ++;
}
```

```
reporter:
while TRUE {
    print_counter;
    counter=0;
}
```

Example – Possible Errors



7. is lost!

Example – Possible Errors

```
counter++    LOAD  ACC, COUNTER  
             INC    ACC  
             SAVE   COUNTER, ACC
```

Race:

- when processes access a shared variable
- when outcome depends on order and running speed of processes
- outcome may be different for different runs

Example – Possible Errors

P1:

```
while TRUE  
  k=k+1;
```

k=0 (initial value)

what about the values of k
depending on the order of
P1 and **P2** executions?

P2:

```
while TRUE  
  k=k+1;
```

SOLUTION: mutual
exclusion

Sharing

- two types of sharing:
 - READ (no need for mutual exclusion)
 - WRITE (mutual exclusion needed)
- for consistency
 - mutual exclusion
 - synchronization

Synchronization

- outcome of programs should not be dependent on running order of processes
- programs working together may need to be synchronized at some points
 - e.g. a program uses output calculated by another program

Mutual Exclusion

critical section (CS): Part of code in a process in which operations on shared resources are performed.

mutual exclusion ensures that only one process executes a CS for a resource at a time

Example

P1:

```
while TRUE {  
    <non-CS>  
    mx_begin  
        <CS ops>  
    mx_end  
    <non-CS>  
}
```


P2:

```
while TRUE {  
    <non-CS>  
    mx_begin  
        <CS ops>  
    mx_end  
    <non-CS>  
}
```

Mutual Exclusion – Possible Problems

- deadlock
 - more than one process requires the same resources
 - each process does not release the resource required by the other

Example: 3 processes and 3 resources



<u>P1()</u>	<u>P2()</u>	<u>P3()</u>
req(R1);	req(R2);	req(R3);
req(R2);	req(R3);	req(R1);

Mutual Exclusion

- `mx_begin`
 - is there a process in its CS which has not yet executed `mx_end` ?
 - if NOT
 - allow process to proceed into CS
 - leave mark for other processes
- `mx_end`
 - allow any process waiting to go into its CS to proceed
 - if there aren't any, then leave mark (empty)

Criteria for Mutual Exclusion Implementation

- only one process may be in its CS
- if a process wants to enter its CS and if there are no others executing their CS, it shouldn't wait
- any process not executing its CS should not prevent another process from entering its own CS
- no assumptions should be made about the order and speed of execution of processes
- no process should stay in its CS indefinitely
- no process should wait to enter its own CS indefinitely

Mutual Exclusion Solutions

- software based solutions
- hardware based solutions
- software and hardware based solutions

A Software Based Solution

- use a shared flag that shows whether a process is in its CS or not: *busy*
 $busy \leftarrow \text{TRUE}$: process in CS
 $busy \leftarrow \text{FALSE}$: no process in CS
- **mx_begin:** `while (busy);`
 `busy = TRUE;`
 - wait until process in CS is finished
 - enter CS
- **mx_end:** `busy = FALSE;`

A Software Based Solution

- a possible error
 - *busy* is also a shared variable!
 - Example:
 - P1 checks and finds *busy*=FALSE
 - P1 interrupted
 - P2 checks and finds *busy*=FALSE
 - both P1 and P2 enter CS

Solutions Requiring *Busy Waiting*

shared variable turn = 1;

Process 1:

local variables

my_turn=1;

others_turn=2;

Process 2:

local variables

my_turn=2;

others_turn=1;

mx_begin: while (turn != my_turn);

mx_end : turn = others_turn;

Solutions Requiring *Busy Waiting*

- uses up CPU time
- works properly but has limitations:
 - processes enter their CS in turn
 - depends on speed of process execution
 - depends on number of processes

Solutions Requiring *Busy Waiting*

- first correct solution: Dekker algorithm
- Peterson algorithm (1981)
 - similar approach to Dekker's algorithm
 - but is simpler

Peterson Algorithm

- shared variables:

req_1, req_2: bool and initialized to FALSE

turn: integer and initialized to "P1" or "P2"

P1:

```
mx_begin:
req_1 = TRUE;
turn = P2;
while (req_2 && turn==P2);
    < CS >
mx_end: req_1 = FALSE;
```

P2:

```
mx_begin:
req_2 = TRUE;
turn = P1;
while (req_1 && turn==P1);
    < CS >
mx_end: req_2 = FALSE;
```

Peterson Algorithm

- different scenarios:
 - P1 is active, P2 is passive
req_1=TRUE and turn=P2
req_2=FALSE so P1 proceeds after while loop
 - P1 in CS, P2 wants to enter CS
req_2=TRUE and turn=P1;
req_1=TRUE so P2 waits in while loop
P2 continues after P1 executes max_end

Peterson Algorithm

- (*different scenarios cntd.*):
 - P1 and P2 want to enter CS at the same time

P1 :

```
req_1=TRUE;
turn=P2;
```

P2 :

```
req_2=TRUE;
turn=P1;
```

⇒ order depends on which process assigns value to the `turn` variable first.

Hardware Based Solutions

- with uninterruptable machine code instructions completed in one machine cycle
 - e.g.: `test_and_set`
 - busy waiting used
 - when a process exits CS, no mechanism to determine which other process enters next
 - indefinite waiting possible
- disabling interrupts
 - interferes with scheduling algorithm of operating system

Hardware Based Solutions

```
test_and_set(a):  cc ← a  
                  a ← TRUE
```

- with one machine instruction, contents of “a” copied into condition code register and “a” is assigned TRUE

```
mx_begin: test_and_set(busy);  
          while (cc) {  
            test_and_set(busy);  
          }  
mx_end:   busy=FALSE;
```

busy: shared variable
cc: local condition code

Semaphores

- hardware and software based solution
- no busy waiting
- does not waste CPU time
- **semaphore** is a special variable
 - only access through using two special operations
 - special operations cannot be interrupted
 - operating system carries out special operations

Semaphores

- s : semaphore variable
- special operations:
 - P (wait): when entering CS: `mutex_begin`
 - V (signal): when leaving CS: `mutex_end`

$P(s)$:

```
if (s > 0)
```

```
    s=s-1;
```

```
else
```

```
    wait_on_s;
```

$V(s)$:

```
if (anyone_waiting_on_s)
```

```
    activate_next_in_line;
```

```
else
```

```
    s=s+1;
```


Semaphores

- take on integer values (≥ 0)
- created through a special system call
- assigned an initial value
- binary semaphore:
 - can be 0/1
 - used for CS
- counting semaphore:
 - can be integers ≥ 0

Example: Observer – Reporter

global shared variables:

counter: integer;

sem: semaphore;

process P1:

observe;

P(sem);

counter++;

V(sem);

....

process P2:

...

P(sem);

print(counter);

counter=0;

V(sem);

main_program:

sem=1; counter=0;

activate(P1);

activate(P2);

Example: Observer – Reporter

sample run:

P1: P(sem) ... sem=0;

P2: P(sem) ... sem=0 so P2 is suspended

P1: V(sem) ... P2 is waiting for sem; activate P2

P2: V(sem) ... no one waiting; sem=1

Synchronization with Semaphores

- a process may require an event to proceed – process is suspended
 - e.g. process waiting for input
 - another process detecting the occurrence of event wakes up suspended process
- ⇒ “suspend – wake-up” synchronization

Synchronization with Semaphores

- solution:

```
event:semaphore; event=0;
```

```
process P1:
```

```
...
```

```
P(event);
```

```
...
```

```
...
```

```
process P2:
```

```
...
```

```
...
```

```
V(event);
```

- more than two processes may be synchronized

Semaphores

Initial value for semaphore:

- =1 for mutual exclusion
- =0 for synchronization

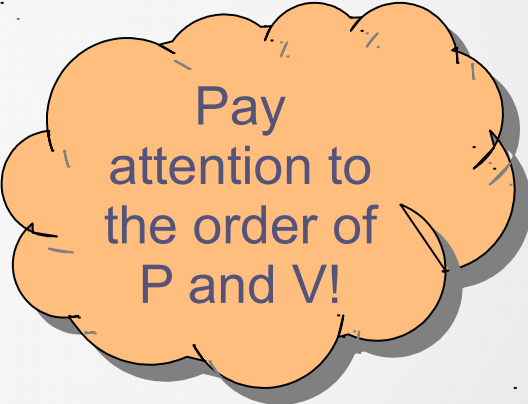
Semaphores

- possible deadlock scenario:

```
x, y: semaphore;  
x=1; y=1;
```

process 1: **process 2:**

...	...
P (x) ;	P (y) ;
...	...
P (y) ;	P (x) ;
...	...
V (x) ;	V (y) ;
V (y) ;	V (x) ;
...	...



Pay
attention to
the order of
P and V!